

---

# Data Processing Notes

---

## CONTENTS

|          |   |          |
|----------|---|----------|
| <b>1</b> | <b>Overview</b>                         | <b>1</b> |
| <b>2</b> | <b>Clamping</b>                         | <b>1</b> |
| <b>3</b> | <b>Imputation</b>                       | <b>2</b> |
| <b>4</b> | <b>Resize</b>                           | <b>2</b> |
| 4.1      | Motivation . . . . .                    | 2        |
| 4.2      | Goals . . . . .                         | 2        |
| 4.3      | Algorithm Statement . . . . .           | 2        |
| 4.4      | Functional Privacy Parameters . . . . . | 3        |
| 4.5      | Examples . . . . .                      | 4        |
| 4.6      | How to choose $n$ . . . . .             | 5        |
| 4.7      | Accuracy/Error . . . . .                | 5        |

## 1 OVERVIEW

This document contains notes on our choices regarding “data processing”, which I am using to stand in for all things related to data bookkeeping (data bounds, imputation, sample size calculations, etc.).

## 2 CLAMPING

Many operations in the library require the definition of a set of possible data values we would like to use for said operation. For numeric variables, this typically takes the form of a closed interval  $[min, max] \in \mathbb{R}$  or  $\mathbb{Z}$ . For categorical variables, this is a discrete set of elements. In our system, we have slightly different notions of clamping for numeric and categorical variables. We consider *f64* and *i64* to be allowable numeric types, and *i64*, *bool*, and *String* to be allowable categorical types.

Our method for clamping numeric types should be familiar; it accepts data, a min, and a max, and maps data elements outside of the interval  $[min, max]$  to the nearest point in the interval. Null values (represented as *NAN* for the *f64* type and which do not exist for the *i64* type) are preserved by numeric clamping.

Our categorical clamping accepts data, a set of feasible values, and a null value. Data elements that are outside the set of feasible values are mapped to the null value. Null values are preserved by the categorical clamping.

[[ Christian: Give algorithm statement here? ]]

### 3 IMPUTATION

Much like clamping, we have slightly different notions of imputation for numeric and categorical variables. In this case, we consider only *f64* to be an allowable numeric type, while *f64*, *i64*, *bool*, and *String* are allowable categorical types.

Numeric imputation is parameterized by min, max, and a supported probability distribution (e.g. Gaussian or Uniform) with appropriate arguments. *NAN* values are replaced with a draw from the provided distribution.

Categorical imputation is parameterized by categories, probabilities, and a null value. Each data element equal to the null value is replaced with a draw from the set of categories, according to the probabilities provided.

[[ Christian: Give algorithm statement here? ]]

### 4 RESIZE

#### 4.1 Motivation

We want our library to support cases both in which the analyst does and does not have access to the number of records in the data. Doing the latter well is an under-explored problem, so we propose the *resize* framework as a possible solution.

#### 4.2 Goals

The *resize* framework is a means of jointly achieving a few different goals:

1. guarantee known  $n$  for analyses that require it,
2. give users flexibility in how they trade off between bias and variance, and
3. allow for both c-stability and privacy amplification from subsampling within the same privacy calculus.

#### 4.3 Algorithm Statement

The *resize* function takes the following inputs:

1.  $X$ : The private underlying data.
2.  $\tilde{n}$ : The size of the private underlying data.
3.  $n$ : The desired size of the new data.

4.  $p$ : The proportion of the underlying data that can be used to construct the new data.  $p$  can be  $> 1$ .
5. ...: Various arguments explaining imputation rules (not of interest for this document).

Let  $sample(Y, m)$  be a function that samples  $m$  elements from data set  $Y$  without replacement. Let  $Aug(Y, m, \dots)$  be a function that imputes new elements independent of the data (using imputation parameters given by  $\dots$ ) for a data set  $Y$  until it is of size  $m$ . The algorithm will look something like the following:

---

**Algorithm 1** Resize:  $resize(X, n, p, neighboring, \dots)$

---

```

1:  $c \leftarrow \lceil p \rceil$  ▷ sets c-stability property
2:  $s \leftarrow p/c$  ▷ sets subsampled_proportion property
3:  $X_c \leftarrow \bigcup_{i=1}^c X$  ▷ create new database of size  $c\tilde{n}$ , composed of  $c$  copies of  $X$ 
4: if neighboring == “replace one” then
5:    $m \leftarrow \lfloor s\tilde{n} \rfloor$  ▷ number of records that can be filled using subsampled private data
6: else if neighboring == “add/remove one” then
7:    $m \leftarrow \text{Binomial}(c\tilde{n}, s)$  ▷ number of records that can be filled using subsampled private data
8: end if
9:  $(\epsilon', \delta') \leftarrow \left( \log(1 + s(e^{c\epsilon} - 1)), s \left( \sum_{i=1}^{c-1} e^{i\epsilon} \right) \delta \right)$  ▷ privacy amplification via subsampling
10:  $X' \leftarrow sample(X_c, \max(m, n), neighboring) \cup (Aug(\emptyset, \max(0, n - m), \dots))$ 
11: return  $(X', \epsilon', \delta')$ 

```

---

The  $\epsilon', \delta'$  terms come from first applying the group privacy definition with group size  $c$  to the database  $X_c$  to get  $(c\epsilon, \left( \sum_{i=1}^{c-1} e^{i\epsilon} \right) \delta)$  [Vad17] and then applying privacy amplification by subsampling results from Theorems 8 and 9 of [BBG18]. Note that, for the “replace one” definition, we could be using  $\frac{m}{c\tilde{n}}$  instead of  $s$  in the privacy calculation. Using  $s$  gives us a very slightly worse privacy guarantee (the only difference is the  $\lfloor \cdot \rfloor$  we used to get  $m$ ), but is nice for consistency between the methods and not having to keep track of  $m$  as an extra property.

#### 4.4 Functional Privacy Parameters

We established in section 4.3 that a user asking for an  $(\epsilon, \delta)$ -DP guarantee will get an  $(\epsilon', \delta')$ -DP guarantee with respect to the original private data. What we’d really like, however, is for the user to ask for an  $(\epsilon, \delta)$ -DP guarantee and have the library come up with what we will call a *functional*  $(\epsilon'', \delta'')$  that will ensure  $(\epsilon, \delta)$ -DP on the original data. Any components that operate on the resized data will use the *functional*  $(\epsilon'', \delta'')$  internally instead of the parameters passed by the user.

**Theorem 1.** *A mechanism that respects*

$$\left( \frac{1}{c} \log \left( \frac{e^{\epsilon'} - 1}{s} + 1 \right), \frac{\delta'}{s \left( \sum_{i=1}^{c-1} e^{i\epsilon} \right)} \right) \text{-DP}$$

*on the resized data respects  $(\epsilon, \delta)$ -DP on the true private data.*

*Proof.* We know that an  $(\epsilon, \delta)$ -DP on the resized data corresponds to a

$$(\epsilon', \delta') = \left( \log(1 + s(e^{c\epsilon} - 1)), s \left( \sum_{i=1}^{c-1} e^{i\epsilon} \right) \delta \right)$$

guarantee on the private data. We just need to invert the function to find  $(\epsilon, \delta)$  in terms of  $(\epsilon', \delta')$ .

Let's start with  $\epsilon, \epsilon'$ :

$$\begin{aligned}\epsilon' &= \log(1 + s(e^{c\epsilon} - 1)) \\ e^{\epsilon'} &= 1 + s(e^{c\epsilon} - 1) \\ \frac{e^{\epsilon'} - 1}{s} &= e^{c\epsilon} - 1 \\ \frac{1}{c} \log\left(\frac{e^{\epsilon'} - 1}{s} + 1\right) &= \epsilon.\end{aligned}$$

We carry out a similar calculation for  $\delta, \delta'$ :

$$\begin{aligned}\delta' &= s \left( \sum_{i=1}^{c-1} e^{i\epsilon} \right) \delta \\ \frac{\delta'}{s \left( \sum_{i=1}^{c-1} e^{i\epsilon} \right)} &= \delta.\end{aligned}$$

□

So, in order for a mechanism to respect  $(\epsilon, \delta)$ -DP on the original data, it must respect  $\left(\frac{1}{c} \log\left(\frac{e^{\epsilon}-1}{s} + 1\right), \frac{\delta}{s(\sum_{i=1}^{c-1} e^{i\epsilon})}\right)$ -DP on the resized data.

We now present an mini-algorithm for finding the *functional*  $(\epsilon, \delta)$ .

---

**Algorithm 2** Finding Functional  $(\epsilon_f, \delta_f)$ : `get_func_priv( $p, \epsilon, \delta$ )`

---

- 1:  $c \leftarrow \lceil p \rceil$  ▷ sets c-stability property
  - 2:  $s \leftarrow p/c$  ▷ sets subsampled\_proportion property
  - 3:  $\epsilon_f \leftarrow \frac{1}{c} \log\left(\frac{e^{\epsilon}-1}{s} + 1\right)$
  - 4:  $\delta_f \leftarrow \frac{\delta}{s(\sum_{i=1}^{c-1} e^{i\epsilon})}$
  - 5: **return**  $(\epsilon_f, \delta_f)$
- 

## 4.5 Examples

Let  $X$  be such that  $\tilde{n} = |X| = 100$ . We can look at a few examples of calls to the `resize` function (which will return  $X'$ ) and check the behavior.

1. `resize( $X, 150, 1, \dots$ )`:  $X'$  will be made up of the 100 true elements of  $X$  and 50 imputed values. The functional privacy parameters are identical to the ones the user provides.
2. `resize( $X, 100, 0.75, \dots$ )`:  $X'$  will be made up of 75 true elements of  $X$  and 25 imputed values. The functional privacy parameters will benefit (lower noise) from amplification via subsampling.
3. `resize( $X, 90, 1.5, \dots$ )`:  $X'$  will be a random sample of  $X \cup X$  of size 90. The functional privacy parameters will lead to greater noise than what the user provides, as they have to take into account the new c-stability of 2. This example is illustrative in that it shows that the functional privacy usage is affected only by  $p$  – it has nothing to do with the relative sizes of the  $X$  and  $X'$ .

## 4.6 How to choose $n$

To this point, we have ignored how a user would go about choosing  $n$ . If no trusted public metadata exist, we suggest estimating the size of the private database (e.g. using the Geometric mechanism) and using that release for  $n$ .

## 4.7 Accuracy/Error

The accuracy/error calculations in the library are always relative to the resized data  $X'$ , rather than the true private data  $X$ . We are currently thinking about how best to extend guarantees from  $X'$  to  $X$ .

## REFERENCES

- [BBG18] Borja Balle, Gilles Barthe, and Marco Gaboardi. Privacy amplification by sub-sampling: Tight analyses via couplings and divergences. *CoRR*, abs/1807.01647, 2018.
- [Vad17] Salil Vadhan. The complexity of differential privacy. In *Tutorials on the Foundations of Cryptography*, pages 347–450. Springer, 2017.